Movie Genre Classification Using Transformer-Based Models: A Multi-Label Approach

Leo Chen (yc687)

Kaisen Yao (ky136)

Duke University

Author Note: [The technical component of this project is presented in a GitHub repository[1].]

[1] GitHub repository: https://github.com/kaisenyao/NLP_Final_ClosedAI

Abstract

This project aims to develop a multi-label classification model for predicting movie genres using transformer-based architectures, specifically BERT. The project is divided into two main phases: (1) training and evaluating the model using synthetic data and (2) applying the model to real-world data. The methodology begins with preprocessing the data, including cleaning the datasets, handling missing values, and transforming plot descriptions into concise summaries using the OpenAI API. The synthetic and real-world datasets are then prepared for multi-label classification, where genres are encoded as binary vectors, allowing the model to predict multiple genres simultaneously. A BERT-based transformer architecture is employed, with the [CLS] token serving as the aggregated representation for genre classification. The model is trained using a Binary Cross-Entropy loss function, and various regularization techniques, such as dropout and early stopping, are applied to ensure robust generalization. Model performance is evaluated through both quantitative metrics, including precision, recall, F1-score, and Hamming loss, and qualitative analysis of genre classification accuracy. This methodology provides a comprehensive framework for building a robust and scalable multi-label genre classification system that can handle the complexities of overlapping genres and varied data inputs.

1. Introduction

Movies often span multiple genres, such as Action, Comedy, or Drama, each representing unique aspects of themes, tone, and storytelling. The task of accurately classifying movies into genres based on plot descriptions is both complex and vital for enhancing film organization, recommendation systems, and analytical insights. This complexity arises from the significant overlap between genres and the nuanced, context-dependent language used in plot descriptions.

We want to address this movie genre classification challenge by employing BERT, a transformer-based model. BERT is particularly suited for this task due to its ability to capture intricate contextual relationships in text, enabling it to discern genre-specific linguistic patterns.

The primary objective of this project is to develop a robust model capable of accurately predicting multiple genres for a given movie. Beyond achieving high classification accuracy, this endeavor aims to uncover linguistic patterns associated with various genres, offering deeper insights into narrative structures and storytelling. This project showcases the potential of state-of-the-art natural language processing techniques in solving real-world multi-label classification challenges.

2. Data

2.1. Raw Data

The raw data used in this project was obtained through a program that legally accesses information from Letterboxd.com, a global social network for film discussion and discovery. Letterboxd allows users to keep track of the movies they've watched, share reviews, and create lists of their favorite movies. The movie-related metadata available on Letterboxd is supplied by The Movie Database (TMDb), ensuring reliable and well-maintained information.

The dataset consists of three primary files:

1. **movies.csv** - Contains basic information about films:
    - id: Movie identifier (primary key)
    - name: The name of the film
    - date: Year of release
    - tagline: The film's slogan
    - description: Plot description
    - minute: Duration of the film (in minutes)
    - rating: Average rating of the film
2. **genres.csv** - Contains information about film genres:
    - id: Movie identifier (foreign key)
    - genre: Genre of the film
3. **themes.csv** - Contains information about film themes:
    - id: Movie identifier (foreign key)
    - theme: Themes of the film

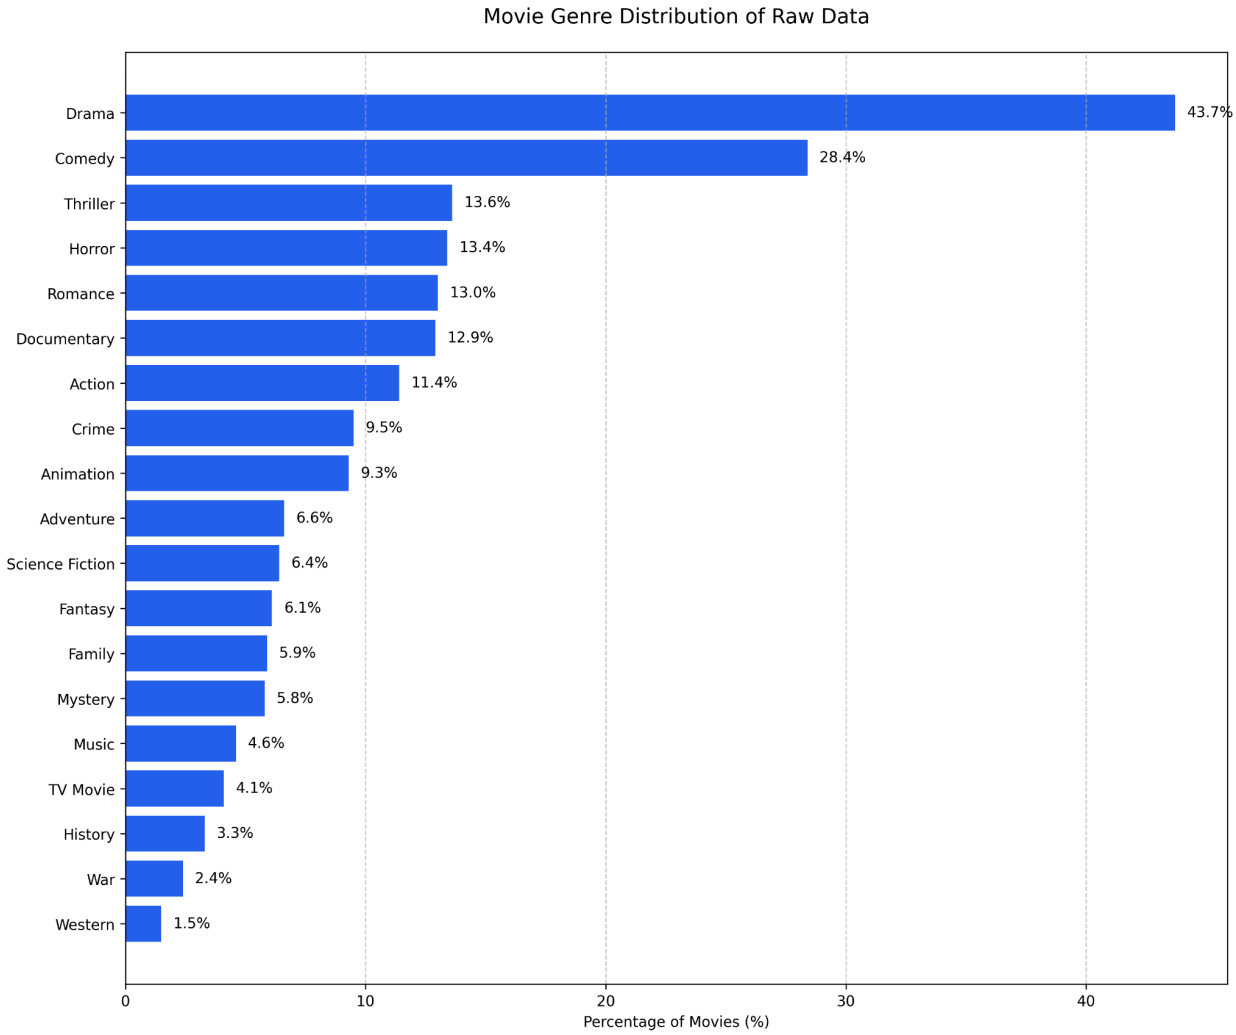

[Figure 1. Movie Genre Distribution of Raw Data]

Based on the analysis of this substantial dataset of 97,554 movies, there are 19 distinct genres, with movies often being categorized under multiple genres, resulting in 197,137 total genre entries. Drama overwhelmingly dominates the dataset, appearing in 43.7% of all movies (42,627 instances), followed by Comedy at 28.4% (27,697 instances). There's a significant drop to the next tier of genres, with Thriller, Horror, Romance, and Documentary each appearing in 12-14% of movies. Action films comprise 11.4% of the dataset, while genres like Crime, Animation, and Adventure range between 6-10%. Notably, some traditional genres appear less frequently: Science Fiction and Fantasy each appear in roughly 6% of movies, while History

(3.3%), War (2.4%), and Western (1.5%) are among the least common genres. These statistics suggest that the film industry heavily favors dramatic and comedic content, while specialized genres like Westerns and War films represent relatively niche categories in the overall landscape of cinema.

2.2. Synthetic Data

Some synthetic data was used when we designed and tried implementing the workflow. To achieve that, we asked ChatGPT to make up 20 pseudo-movies with their corresponding plot descriptions and genre classifications, then store them into file **synthesis_data.csv**. This synthetic dataset was carefully engineered to mimic the structure and genre diversity of real data, providing a testbed for debugging and refining the workflow.

3. Methodology

This project follows a structured methodology consisting of three key stages: data cleaning and manipulation, model training, and evaluation. The process begins with data preprocessing to ensure consistency and quality, including cleaning missing or non-English content, consolidating themes and genres, and preparing the data for multi-label classification. Next, the model, built on a transformer-based architecture leveraging BERT, is trained to handle overlapping and diverse genres using techniques like Binary Cross-Entropy loss minimization and regularization. Finally, the model's performance is evaluated using both synthetic and real-world datasets through quantitative metrics and qualitative analysis, providing insights into strengths and areas for improvement.

3.1. Data Cleaning and Manipulation

Data preprocessing was a crucial step to ensure the consistency and quality of the input data. We began by addressing missing values, duplicate entries, and non-English content, all of which were removed to create a clean and reliable dataset. Then, leveraging the OpenAI API, we transformed the plot descriptions into concise, meaningful summaries, eliminating redundancy while preserving essential narrative elements. This significantly reduced noise and made the data more suitable for model training.

Next, we loaded three datasets: movies.csv, which contained movie descriptions; themes.csv, which provided thematic information; and genres.csv, which listed genre classifications. From movies.csv, we retained only the id and description columns to focus on the

core narrative data. This filtering ensured that irrelevant details were excluded, leaving behind a concise and relevant dataset for downstream tasks.

To prepare the data for multi-label classification, we consolidated themes and genres by grouping them based on the id. Multiple entries for the same movie were combined into single strings, separated by a pipe (|), creating a unified representation of all themes and genres associated with each movie. This step made it easier to handle movies with overlapping labels and ensured that each movie's unique characteristics were captured.

The datasets were then merged on the id column, integrating movie descriptions, themes, and genres into a single cohesive table. Rows with missing values in key columns, such as id, description, theme, or genre, were removed to maintain data integrity and eliminate incomplete entries that could negatively impact the model's performance. This ensured that the dataset was both comprehensive and reliable.

To create a manageable subset for training and evaluation, we randomly sampled 1,000 rows from the cleaned dataset. This sampling was performed with a fixed random seed, ensuring reproducibility while maintaining a representative selection of data. Finally, the processed dataset was saved to a new CSV file, ready for the modeling phase.

This thorough data preprocessing pipeline offered several key benefits. First, it significantly improved the quality and consistency of the dataset by removing irrelevant or incomplete entries. Second, by consolidating themes and genres into a unified format, the pipeline optimized the data for multi-label classification, simplifying the handling of overlapping genres. Finally, transforming plot descriptions into concise summaries reduced noise in the input data, enabling the model to focus on the most important narrative details, ultimately improving

learning efficiency. These steps collectively ensured that the model received structured, high-quality input, providing a strong foundation for subsequent training and evaluation.

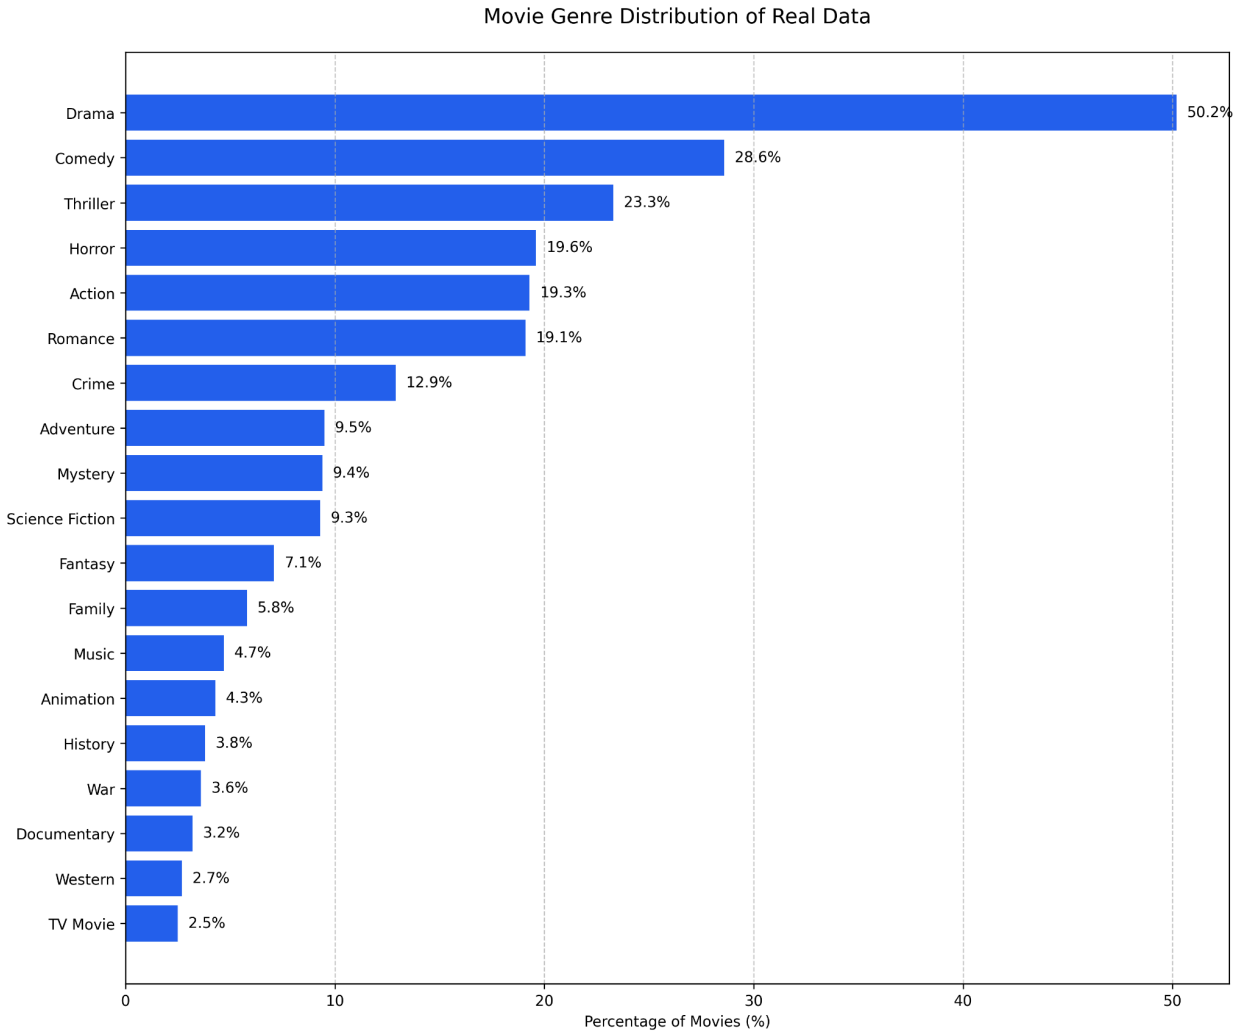

[Figure 2. Movie Genre Distribution of Real Data]

Based on the analysis of the real dataset, there are 19 distinct genres represented. Drama dominates the collection, appearing in just over half (50.2%) of all movies, followed by Comedy at 28.6% and Thriller at 23.3%. Action, Horror, and Romance each appear in roughly one-fifth of the movies (19-20%), showing their significant presence in the dataset. Crime films make up 12.9% of the collection, while Adventure, Mystery, and Science Fiction each appear in approximately 9-10% of movies. Fantasy, Family, Music, and Animation are moderately

represented (between 4-7%). The least common genres are TV Movie (2.5%), Western (2.7%), Documentary (3.2%), War (3.6%), and History (3.8%). This distribution suggests that the dataset predominantly features dramatic content, with strong representation of lighter entertainment (Comedy) and suspense (Thriller), while specialized genres like Westerns and TV Movies appear less frequently. It's worth noting that many movies are tagged with multiple genres, which is why the percentages sum to more than 100%.

3.2. Model Architecture and Training Principles

The model leverages BERT, a state-of-the-art transformer architecture renowned for its ability to capture intricate contextual relationships in text. The input to the model consists of tokenized movie descriptions, which are processed into contextual embeddings for each token in the sequence. A crucial component of this architecture is the [CLS] token, which aggregates information from the entire input sequence and serves as a summary representation for classification tasks.

To adapt BERT for multi-label classification, the [CLS] token embedding is passed through a fully connected classification layer. This layer outputs independent probabilities for each genre, computed using a sigmoid activation function. This approach ensures compatibility with multi-label tasks, treating each genre as an independent binary classification problem. By using the sigmoid activation, the model assigns a probability score to each genre, reflecting its confidence in the presence of that genre for a given movie.

The model was implemented using PyTorch[2], with the AutoTokenizer and AutoModelForSequenceClassification utilities from the Hugging Face Transformers library[3].

[2] Reference to PyTorch: https://pytorch.org/
[3] Reference to Hugging Face Transformers: https://huggingface.co/docs/transformers/

The MovieGenreDataset class facilitates data handling, including tokenization, truncation, padding, and creation of PyTorch tensors for input features and labels. This modular design simplifies data preprocessing and ensures efficient batch processing during training.

For training, the model minimizes the Binary Cross-Entropy loss function, which is particularly suited for multi-label classification as it evaluates each label independently. The synthetic and real-world datasets were split into training and testing subsets to evaluate the model's performance on unseen data. The training pipeline employs iterative optimization using gradient descent, with the Adam optimizer being a common choice for its adaptive learning rate capabilities.

To enhance generalization and prevent overfitting, several regularization techniques were applied. Dropout layers were used to randomly deactivate a subset of neurons during training, reducing reliance on specific pathways. Early stopping was also employed, halting training once validation performance ceased to improve, thus preventing unnecessary overfitting.

Finally, the model architecture is designed to scale efficiently on GPU-accelerated hardware, utilizing batching and parallel computation to handle the high-dimensional input embeddings generated by BERT. This comprehensive approach ensures robust adaptability to diverse and complex data inputs, making the model well-suited for real-world multi-label classification tasks such as movie genre prediction.

3.3. Evaluations

For this multi-label classification problem, precision, recall, F1-score, and Hamming loss are computed to evaluate the performance of the model on both synthetic and real-world

datasets. These metrics were chosen since they address key aspects of multi-label classification: quality of predictions, coverage of true labels, and the overall correctness of the output.

1. **Precision**: measures the proportion of correctly predicted genres among all genres predicted by the model. It is calculated for each genre as:

$$Precision_i = \frac{True\ Positives_i}{True\ Positives_i + False\ Positives_i}$$

where True Positives refer to the number of instances where the model correctly predicts a genre that a movie belongs to, while False Positives denote the instances where the model incorrectly predicts a genre that is not actually associated with the movie.

The average precision is computed across all genres using macro-averaging, which treats all genres equally, regardless of their frequency.

2. **Recall**: Measures how well the model identifies all true genres for each movie. And recall for each genre is computed as:

$$Recall_i = \frac{True\ Positives_i}{True\ Positives_i + False\ Negatives_i}$$

where False Negatives is the number of times the model fails to assign a genre that is true for the movie.

Like precision, recall is averaged across all genres using macro-averaging to ensure that underrepresented genres are not overshadowed by more frequent ones.

3. **F1-Score**: combines precision and recall into a single metric by calculating their harmonic mean:

$$F1 - Score_i = 2 \times \frac{Precision_i \times Recall_i}{Precision_i + Recall_i}$$

This metric is especially useful for evaluating models where there is a trade-off between precision and recall.

4. **Hamming Loss**: quantifies the fraction of incorrect labels (both false positives and false negatives) among all possible labels:

$$Hamming\ Loss = \frac{1}{N \times L} \sum_{i=1}^{N} \sum_{j=1}^{L} 1[y_{ij} \neq \widehat{y_{ij}}]$$

where N is the number of samples, L is the number of labels, y_{ij} is the ground truth label, and $\widehat{y_{ij}}$ is the predicted label. A lower Hamming loss indicates better overall classification performance.

From above, precision and recall are critical metrics for multi-label classification as they separately evaluate prediction accuracy and label coverage. Precision focuses on the proportion of correctly predicted genres among all predicted genres, providing insight into the model's ability to avoid irrelevant or false-positive classifications. On the other hand, recall measures the model's effectiveness in identifying all true genres, ensuring that no relevant genres are missed. This distinction is particularly important in genre classification tasks, where false positives (irrelevant genres) and false negatives (missing genres) can have different consequences. Additionally, macro-averaging ensures that all genres, including underrepresented ones, are given equal importance during evaluation, addressing the inherent imbalance in the dataset.

The F1-score is the harmonic mean of precision and recall, offering a balanced evaluation metric when the model needs to trade off between these two aspects. In multi-label classification, achieving high precision often comes at the cost of recall and vice versa. By combining these two metrics, the F1-score provides a single, interpretable value that reflects the model's overall performance. It is particularly useful in scenarios where neither precision nor recall alone can fully capture the effectiveness of the model.

Hamming loss provides a broader perspective by quantifying both false positives and false negatives across all genres. Unlike precision and recall, which focus primarily on positive predictions, Hamming loss evaluates the fraction of misclassified labels relative to the total number of labels. This metric is especially well-suited for multi-label classification, as it penalizes errors for both predicted and omitted genres. By highlighting overall classification errors, Hamming loss ensures that both incorrect and missed labels are taken into account, making it a robust metric for assessing multi-label tasks.

4. Outcomes

4.1. Results

After evaluating the model on both synthetic and real-world datasets, the performance metrics are summarized in the table below:

Dataset	Precision	Recall	F1-Score	Hamming Loss
Synthetic	0.136	0.273	0.164	0.308
Real-World	0.368	0.148	0.168	0.122

Synthetic data was initially used to validate the pipeline in a controlled environment. After preprocessing and tokenization, the model was trained to predict genre labels. Evaluation on the synthetic dataset revealed moderate recall (0.273) but low precision (0.136), resulting in an F1-score of 0.164 and a Hamming loss of 0.308. These metrics indicated that while the model could detect multiple genres, it frequently misclassified or overpredicted certain genres. For example, a description tagged as “Adventure” and “Comedy” was sometimes misclassified as “Action” or “Music.” Qualitative analysis suggested that the errors stemmed from the limited representation of rare genres in the synthetic dataset, prompting adjustments to the data generation process to improve balance and representation.

In the second phase, the model was applied to a legally acquired real-world dataset. Preprocessing steps were aligned with those used for the synthetic data to ensure consistency and facilitate a seamless transition. The same transformer-based model architecture was employed,

leveraging its capacity to generalize contextual relationships learned from synthetic to real data. Evaluation on the real dataset showed improved performance, with precision increasing to 0.368 and Hamming loss decreasing to 0.122. This improvement highlighted the model's ability to adapt to more complex and diverse inputs. It excelled at predicting common genres like "Drama" and "Comedy" but continued to struggle with rare or ambiguous genres, such as "Fantasy" or "Thriller." For instance, a movie tagged as "Comedy" was occasionally misclassified as "Thriller" due to overlapping or noisy textual descriptions.

An F1-score of 0.168 across phases underscores the model's suboptimal performance, as this metric represents the harmonic mean of precision and recall, where a perfect balance is scored at 1. The low F1-score suggests the model struggles to balance precision (correctness of positive predictions) and recall (coverage of actual positives). In this case, the model either predicts very few true positives or generates many false positives, likely due to class imbalances within the dataset. The bias towards majority classes further marginalizes minority classes, exacerbating the issue. To improve these results, strategies such as hyperparameter tuning, incorporating class-specific weighting, or resampling techniques like oversampling rare genres or undersampling dominant ones should be explored.

4.2. Insights from Model Performance

The transformer architecture's strength lies in its ability to capture semantic relationships between words, enabling accurate genre classification in many cases. The use of the [CLS] token provided a holistic view of the input sequence, allowing the model to make nuanced predictions. However, the results also revealed limitations, such as difficulty distinguishing overlapping genres and handling underrepresented classes.

Evaluation results highlight both the strengths and weaknesses of the model. Improved precision on the real-world dataset demonstrates the model's ability to effectively learn from more complex and diverse inputs. This indicates that the model performs well when provided with richer and more representative data, leading to fewer false-positive predictions. However, this improvement in precision does not fully compensate for the challenges posed by low recall, which consistently remained low across both synthetic and real-world datasets. This suggests that the model struggles to capture all true genres, particularly rare or underrepresented ones.

Additionally, differences in Hamming loss between the synthetic and real-world datasets emphasize the influence of dataset size and diversity on model performance. The lower Hamming loss for the real-world data indicates that the larger, more diverse dataset allowed the model to make fewer overall errors compared to the synthetic phase. While the real-world dataset enabled the model to better capture genre-specific relationships, it also highlighted challenges in handling rare or ambiguous genres, such as "Fantasy" or "Thriller." For example, descriptions with overlapping or noisy plot elements were often misclassified, underscoring the need for further refinement.

By combining synthetic and real data, the methodology facilitated iterative refinement of the pipeline. The synthetic phase helped identify and address foundational issues, such as label imbalances and overprediction tendencies, while the real data phase validated the model's practical application in a more realistic setting. Quantitative metrics provided a clear view of the model's strengths and weaknesses, while qualitative analysis offered deeper insights into error patterns.

These findings underscore the importance of using robust datasets that capture the full complexity of the task. Potential strategies for future improvement include augmenting the

training data to ensure better coverage of rare genres, balancing genre representation to reduce bias toward frequent labels, and refining the model architecture to better handle overlapping genres. Addressing these challenges will enhance the model's overall performance and increase its applicability to real-world scenarios.

5. Discussion

5.1. Strengths in Quality and Correctness

The transformer-based model excels in contextual understanding, particularly in identifying common genres such as “Drama” and “Comedy.” This strength is further supported by its ability to generalize effectively across both synthetic and real-world datasets. The model demonstrates significant improvements in precision for frequently occurring genres when applied to real-world data, showcasing its adaptability to complex scenarios with high-quality inputs. Metrics such as reduced Hamming loss on real-world datasets underscore its effectiveness in handling multi-label classification tasks, making it a robust solution for genre prediction.

5.2. Challenges with Rare and Overlapping Genres

Despite its strengths, the model faces notable challenges with rare or overlapping genres. Distinguishing between categories such as “Fantasy” and “Science Fiction” or accurately identifying underrepresented genres like “Documentary” proves difficult. Misclassifications often arise from descriptions with ambiguous or nuanced language, leading to lower recall for these less common genres. These limitations suggest that additional training data or improved balancing of genre representation could significantly enhance the model’s correctness.

5.3. Resource and Computational Demands

The use of a pre-trained transformer architecture reduces the need for extensive training from scratch, leveraging transfer learning to achieve reasonable performance with limited labeled

data. Tokenization ensures consistent input lengths, which optimizes computational efficiency during both training and inference. However, the computational demands of fine-tuning transformers are substantial, requiring high-performance hardware such as GPUs or TPUs. This increases resource costs and training time, particularly for large datasets. Moreover, the preprocessing pipeline, including the condensing of descriptions into concise elements, adds to the overall time and resource consumption.

5.4. Mixed Interpretability

The model's interpretability is a double-edged sword. On the one hand, its multi-label binary classification framework provides clear outputs, with genre probabilities offering straightforward insights into predictions. The attention mechanisms in the transformer architecture can also be visualized to understand which parts of a description most influence the model's decisions. On the other hand, the model functions as a black-box system, making it challenging to fully interpret the reasoning behind specific misclassifications. This is particularly evident in cases with abstract or subtle descriptions, where external context or manual validation may be required, reducing practical interpretability.

5.5. Generalizability and Flexibility

The model demonstrates impressive generalizability, transitioning effectively from synthetic to real-world datasets. Its ability to handle multi-label classification ensures robustness in scenarios where movies belong to multiple genres, reflecting the complexities of real-world data. However, its performance is highly dependent on the quality and balance of the training dataset. Insufficient representation of rare genres or noisy descriptions can lead to skewed

predictions and reduced generalizability. Adapting the model to new domains or datasets with different structures would require significant retraining and preprocessing adjustments, posing additional challenges.

In summary, the model combines the contextual strengths of transformer architectures with a robust classification framework, achieving strong performance for common genres and demonstrating adaptability across diverse datasets. However, its high computational requirements, challenges with rare genres, and limited interpretability in edge cases highlight areas for improvement. These limitations call for further optimization to balance performance, efficiency, and interpretability, ensuring the model's reliability and scalability in real-world genre classification tasks.

6. Conclusion

In this project, we successfully developed a multi-label classification model for movie genre prediction using a transformer-based architecture, specifically BERT. Through two key phases—training on synthetic data and applying the model to real-world data—we demonstrated the model's capability to handle complex and diverse inputs. The preprocessing pipeline, including data cleaning, transformation of plot descriptions, and encoding of genre labels, ensured the data was suitable for model training, reducing noise and improving the quality of the input data. The use of the BERT model allowed us to capture nuanced semantic relationships between words, leading to improved accuracy in genre classification.

Despite the improvements in precision with real-world data, the model faced challenges in recall, particularly with rare or ambiguous genres. The evaluation metrics, including precision, recall, F1-score, and Hamming loss, highlighted areas for further optimization, such as better handling of underrepresented genres and fine-tuning the model for overlapping labels. Future work will focus on expanding the dataset, balancing genre representation, and exploring advanced techniques for boosting recall without compromising precision.

Overall, this project provides a solid foundation for multi-label classification tasks in natural language processing. By leveraging transformer models like BERT, we have created a flexible and scalable framework for predicting multiple genres in movie datasets. This approach can be extended to other domains that require multi-label classification, such as document categorization or product tagging.